

In this project, I will clean up a database and perform exploratory data analysis on it. I will do this by

- removing duplicates
- standardizing data
- addressing null values or blanks
- removing any rows and columns not needed
- conduct an EDA

I start off by looking at a snapshot of our data. After first impressions, I create a duplicate of the database and store the original away to be referred to whenever needed.

```
1 • USE sql_project;
2
3   /*looking at a snapshot of our database*/
4 • SELECT *
5   FROM layoffs;
6
7   /*creating a duplicate of the dataset*/
8 • CREATE TABLE layoffs_duplicate
9   LIKE layoffs;
10
11  /*checking that the table has been created with the correct columns*/
12 • SELECT *
13   FROM layoffs_duplicate;
```

layoffs 1 x

	company	location	industry	total_laid_off	percentage_laid_off	date	stage	country	funds_raised_millions
▶	Atlassian	Sydney	Other	500	0.05	3/6/2023	Post-IPO	Australia	210
	SiriusXM	New York City	Media	475	0.08	3/6/2023	Post-IPO	United States	525
	Alerzo	Ibadan	Retail	400	NULL	3/6/2023	Series B	Nigeria	16
	UpGrad	Mumbai	Education	120	NULL	3/6/2023	Unknown	India	631
	Loft	Sao Paulo	Real Estate	340	0.15	3/3/2023	Unknown	Brazil	788
	Embark Trucks	SF Bay Area	Transportation	230	0.7	3/3/2023	Post-IPO	United States	317
	Lendi	Sydney	Real Estate	100	NULL	3/3/2023	Unknown	Australia	59
	Umoja	SF Bay Area	Mediation	62	NULL	3/3/2023	Acquired	United States	152

Using row_number() over(), I am able to suss out where a particular row has repeat datapoints in all columns meaning it is a duplicate.

```

25  /*checking for duplicates within the dataset using row number over*/
26  • SELECT *, ROW_NUMBER() OVER(
27    PARTITION BY company,location,industry,total_laid_off,percentage_laid_off,'date',stage,country,funds_raised_millions) AS row_num
28    FROM layoffs_duplicate;
29
30  /*creating a cte to filter for duplicates*/
31  • WITH duplicate_cte AS (
32    SELECT *, ROW_NUMBER() OVER(
33      PARTITION BY company,location,industry,total_laid_off,percentage_laid_off,'date',stage,country,funds_raised_millions) AS row_num
34    FROM layoffs_duplicate
35  )
36  SELECT *
37  FROM duplicate_cte
38  WHERE row_num >1;
39
40  /*a row number above 1 depicts a duplicate value*/

```

Result Grid

	company	location	industry	total_laid_off	percentage_laid_off	date	stage	country	funds_raised_millions	row_num
▶	Better.com	New York City	Real Estate	NULL	NULL	8/26/2022	Unknown	United States	905	2
	Casper	New York City	Retail	NULL	NULL	9/14/2021	Post-IPO	United States	339	2
	Cazoo	London	Transportation	750	0.15	6/7/2022	Post-IPO	United Kingdom	2000	2
	Elemy	SF Bay Area	Healthcare	NULL	NULL	12/5/2022	Series B	United States	323	2
	ExtraHop	Seattle	Security	NULL	NULL	4/22/2020	Series C	United States	61	2

I create a second duplicate table and populate it with the unique data points, removing the duplicates.

```

52  /*creating a new table that contains only unique rows*/
53  • CREATE TABLE layoffs_dup2(
54    company TEXT,
55    location TEXT,
56    industry TEXT,
57    total_laid_off INT,
58    percentage_laid_off TEXT,
59    date TEXT,
60    stage TEXT,
61    country TEXT,
62    funds_raised_millions INT,
63    row_num INT
64  );
65

```

```

65
66 • INSERT layoffs_dup2
67 WITH duplicate_cte AS (
68 SELECT *, ROW_NUMBER() OVER(
69 PARTITION BY company,location,industry,total_laid_off,percentage_laid_off,'date',stage,country,funds_raised_millions) AS row_num
70 FROM layoffs_duplicate
71 )
72 SELECT *
73 FROM duplicate_cte
74 WHERE row_num = 1;
75
76 • SELECT *
77 FROM layoffs_dup2;
78
79 /*confirming that all rows have a row_num of 1*/
80 • SELECT DISTINCT row_num

```

company	location	industry	total_laid_off	percentage_laid_off	date	stage	country	funds_raised_millions
Included Health	SF Bay Area	Healthcare	NULL	0.06	2022-07-25	Series E	United States	272
8Open	Dublin	Marketing	9	0.09	2022-11-17	Series A	Ireland	35
#Paid	Toronto	Marketing	19	0.17	2023-01-27	Series B	Canada	21
100 Thieves	Los Angeles	Consumer	12	0.03	2022-07-13	Series C	United States	120
10X Genomics	SF Bay Area	Healthcare	100	0.08	2022-08-04	Post-IPO	United States	242

layoffs_dup23 x

We proceed using the second duplicate table as our working dataset.

On to standardization, the column “company” contains words and spaces. The spaces sometimes come before the words or after, creating inconsistencies that may affect future analysis. I us the trim function to solve this issue.

```

73 FROM duplicate_cte
74 WHERE row_num = 1;
75
76 • SELECT *
77 FROM layoffs_dup2;
78
79 /*confirming that all rows have a row_num of 1*/
80 • SELECT DISTINCT row_num
81 FROM layoffs_dup2;
82
83 /*standardizing the data*/
84
85 /*removing any white spaces surrounding company*/
86 • UPDATE layoffs_dup2
87 SET company = TRIM(company);
88
89 • SELECT DISTINCT industry
90 FROM layoffs_dup2
91 ORDER BY 1;
92
93 /*we notice a few issues in the industry column to be addressed
94 -crypto / crypto currency
95 -null values
96 -blanks*/

```

I move to the next column “industry”. The industry column has entries under both “Crypto” and “Crypto Currency” which I decide are the same thing. Since “Crypto” appears more frequently, I would be using that. Industry also has some null values and blanks.

```

88
89 • SELECT DISTINCT industry
90 FROM layoffs_dup2
91 ORDER BY 1;
92
93 /*we notice a few issues in the industry column to be addressed
94 -crypto / crypto currency
95 -null values
96 -blanks*/
97
98 • SELECT *
99 FROM layoffs_dup2
100 WHERE industry LIKE 'Crypto%';
101
102 • UPDATE layoffs_dup2
103 SET industry = 'Crypto'
104 WHERE industry LIKE 'Crypto%';
105
106 • SELECT DISTINCT location
107 FROM layoffs_dup2
108 ORDER BY 1;
109
110 • SELECT DISTINCT country
111 FROM layoffs_dup2

```

Next, I look at where a company in the same location appears more than once but has industry nulls or blanks in one of the entries. To make things easier, I will set the blanks to null values, then populate the nulls with industry values inferred from the existing data.

```

/*addressing null values*/
/*where do we see the same company in the same location having industry null and not null*/
• SELECT *
  FROM layoffs_dup2 AS l1
 JOIN layoffs_dup2 AS l2
   USING (company)
 WHERE (l1.industry IS NULL OR l1.industry = '')
    AND l2.industry IS NOT NULL;

/*setting blank spaces to null*/
• UPDATE layoffs_dup2
  SET industry = NULL
 WHERE industry = '';

/*inputing inferred value for industry column from existing data*/
• UPDATE layoffs_dup2 AS l1
 JOIN layoffs_dup2 AS l2
   USING (company)
  SET l1.industry = l2.industry
 WHERE l1.industry IS NULL
    AND l2.industry IS NOT NULL;

/*checking that inferable industry null values have been populated*/
• SELECT *

```

After confirming that inferable entries have been made, I move to the next column “country”. There are 2 different “United States” because one has a period at the end. We need to remove said period so that the system recognizes them as one.

```

128      /*checking that inferable industry null values have been populated*/
129 •    SELECT *
130      FROM layoffs_dup2
131      WHERE industry IS NULL;
132
133 •    SELECT DISTINCT location
134      FROM layoffs_dup2
135      ORDER BY 1;
136
137 •    SELECT DISTINCT country
138      FROM layoffs_dup2
139      ORDER BY 1;
140      /*united states is showing a little discrepancy*/
141 •    SELECT *
142      FROM layoffs_dup2
143      WHERE country LIKE 'United States%';
144
145 •    UPDATE layoffs_dup2
146      SET country = TRIM(TRAILING '.' FROM country)
147      WHERE country LIKE 'United States%';
148
149 •    SELECT DISTINCT country
150      FROM layoffs_dup2
151      ORDER BY 1;

```

Checking the data types of the columns, I discover that date column is a text data type. I convert it to the more appropriate 'date' data type.

```

152
153      /*checking the datatypes of our columns*/
154 •    DESCRIBE layoffs_dup2;
155
156      /*converting the date to date format*/
157 •    UPDATE layoffs_dup2
158      SET date = STR_TO_DATE(date, '%m/%d/%Y');
159
160      /*now that it's in a date format, we can easily convert to date datatype*/
161 •    ALTER TABLE layoffs_dup2
162      MODIFY COLUMN date DATE;
163

```

My analysis will focus on the numeric aspect of the dataset, having null values the numeric columns will be problematic.

- `SELECT *`
`FROM layoffs_dup2`
`WHERE total_laid_off IS NULL`
`AND percentage_laid_off IS NULL;`
 /*348 rows have both these key columns null, roughly 15% of our dataset. seeing no viable way to populate these cells and yet realising not removing them will heavily influence my analysis moving forward. I opt to delete them. I can always fall back on the original database if need be*/
- `DELETE`
`FROM layoffs_dup2`
`WHERE total_laid_off IS NULL`
`AND percentage_laid_off IS NULL;`

`/*row_num column is now redundant*/`
- `ALTER TABLE layoffs_dup2`
`DROP COLUMN row_num;`

Checking for null values in our 2 numeric columns, I find that 348 rows have both numeric columns null, this accounts for about 15% of our original data. Unfortunately, I can not infer these values from the existing data, nor can I confidently input a calculated estimate. I decided that leaving these rows in will affect my analysis so opt to remove them. Also, the row_num column I used to find the duplicates is no longer needed so I remove that as well.

Feeling content with the cleaning for now, I move on to my exploratory data analysis. Things I am curious about the data set include:

- What is the timeframe of the data set ?
- Which companies went completely under during the period ?
- Which companies are laying off the most employees ?
- Which industry laid off the most on average ?
- Which countries laid off the most ?
- How did layoffs progress over the period of interest ?
- How the stage a company is in affected layoffs.

- ```
/*what date range are we working with in this database*/
SELECT MIN(date),MAX(date)
FROM layoffs_dup2;
```
- ```
/*which companies went completely under */  
SELECT *  
FROM layoffs_dup2  
WHERE percentage_laid_off = 1  
ORDER BY funds_raised_millions DESC;
```
- ```
/*which companies are laying off a large number of employees*/
SELECT company,SUM(total_laid_off)
FROM layoffs_dup2
GROUP BY company
ORDER BY 2 DESC;
```
- ```
/*what industry got hit the most*/  
SELECT industry,AVG(total_laid_off)  
FROM layoffs_dup2  
GROUP BY industry  
ORDER BY 2 DESC;
```

The data set is from 11th March 2020 to 6th March 2023, starting at the peak of the world wide outbreak of the corona virus and afterwards.


```

187
188  /*which companies went completely under */
189 • SELECT industry, COUNT(company)
190 FROM layoffs_dup2
191 WHERE percentage_laid_off = 1
192 GROUP BY industry
193 ORDER BY 2 DESC;

```

Result Grid			Filter Rows:	Export:	Wrap Cell Content:
	industry	COUNT(company)			
▶	Retail	13			
	Food	13			
	Finance	12			
	Education	9			
	Healthcare	7			
	Crypto	7			
	Transportation	6			
	Real Estate	6			
	Marketing	5			
	Media	5			
	Consumer	5			
	Travel	4			
	Other	4			
	Infrastructure	3			
	Product	3			
	Recruiting	2			
	Support	2			
	Fitness	2			

At the top of the list of industries we have retail, food and finance.

```

194
195  /*which companies are laying off a large number of employees*/
196 • SELECT company,SUM(total_laid_off)
197 FROM layoffs_dup2
198 GROUP BY company
199 ORDER BY 2 DESC;

```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:	Fetch rows:
company	SUM(total_laid_off)				
Amazon	18150				
Google	12000				
Meta	11000				
Salesforce	10090				
Microsoft	10000				
Philips	10000				
Ericsson	8500				
Uber	7585				
Dell	6650				
Booking.com	4601				
Cisco	4100				
Peloton	4084				
Byju's	4000				
Carvana	4000				
Twitter	3940				
Better.com	3900				
IBM	3900				
Groupon	3800				
Bvtedance	3750				

Large companies like Amazon, Google and Meta laid off the most employees during the period. However, these companies also have a disproportionately large number of employees in general.

```

200
201      /*what industry got hit the most*/
202 •   SELECT industry,COUNT(company) ,SUM(total_laid_off)
203      FROM layoffs_dup2
204      GROUP BY industry
205      ORDER BY 3 DESC;
206
207      /*what countries got hit the most*/

```

Result Grid Filter Rows: <input type="text"/> Export: Wrap Cell Content:			
	industry	COUNT(company)	SUM(total_laid_off)
▶	Consumer	101	45182
	Retail	163	43613
	Other	106	36209
	Transportation	128	33548
	Finance	239	28344
	Healthcare	163	25894
	Food	115	22855
	Real Estate	100	17565
	Travel	56	17159
	Hardware	13	13828
	Education	79	13338
	Sales	34	13216
	Crypto	89	10693
	Marketing	123	10258
	Fitness	26	8748
	Security	62	5979
	Infrastructure	32	5785

Consumer and retail are leading industries in layoffs.

```

206
207      /*what countries got hit the most*/
208 •   SELECT country,COUNT(country),SUM(total_laid_off)
209      FROM layoffs_dup2
210      GROUP BY country

```

Result Grid Filter Rows: <input type="text"/> Export: Wrap Cell Content:			
	country	COUNT(country)	SUM(total_laid_off)
▶	United States	1291	256420
▶	India	139	35793
	Canada	90	6319
	Brazil	69	10391
	Germany	62	8701
	United Kingdom	58	6398
	Israel	47	3638
	Australia	46	2324
	Singapore	26	5995
	Indonesia	20	3521
	Sweden	17	11264
	China	13	5905
	Netherlands	12	17220
	Nigeria	11	1882
	United Arab Emirates	6	995
	Kenya	6	349
	France	5	915
	Argentina	5	323
	Estonia	5	333
	Ireland	4	257

USA contributing the most companies to the database naturally also has the most number of layoffs.

```
214 • SELECT YEAR(date),COUNT(DISTINCT(MONTH(date))),SUM(total_laid_off)
215 FROM layoffs_dup2
216 WHERE date IS NOT NULL
217 GROUP BY YEAR(date)
218 ORDER BY 1;
```

	YEAR(date)	COUNT(DISTINCT(MONTH(date)))	SUM(total_laid_off)
▶	2020	10	80998
	2021	11	15823
	2022	12	160322
	2023	3	125677

The layoffs were lowest in 2020, but subsequent years saw higher numbers in the aftermath of the pandemic. Though 2023 shows a reduction in layoffs, only 3 months out of the year are accounted for.

```
232 /*over the period, which companies have layed the most employees off*/
233 • WITH company_year(company, year,total_layoffs) AS
234 (
235 SELECT company, YEAR(date),SUM(total_laid_off)
236 FROM layoffs_dup2
237 GROUP BY company,YEAR(date)
238 ORDER BY company ASC),
239 company_year_rank AS (SELECT *,
240 DENSE_RANK() OVER(PARTITION BY year ORDER BY total_layoffs DESC) AS ranking
241 FROM company_year
242 WHERE year IS NOT NULL)
243 SELECT *
244 FROM company_year_rank
245 WHERE ranking <= 5
246 ORDER BY year;
```

	company	year	total_layoffs	ranking
▶	Uber	2020	7525	1
	Booking.com	2020	4375	2
	Groupon	2020	2800	3
	Swiggy	2020	2250	4
	Airbnb	2020	1900	5

Result 17 ×

The top 5 companies laying off employees over the period was quite volatile.

```

248      /*does the stage a company is in affect layoffs*/
249 •    SELECT stage,SUM(total_laid_off)
250      FROM layoffs_dup2
251      GROUP BY stage
252      ORDER BY 2 DESC;

```

Result Grid		
Filter Rows: <input type="text"/>		
Export:		
Wrap Cell Content:		
stage	SUM(total_laid_off)	
Post-IPO	204073	
Unknown	40716	
Acquired	27496	
Series C	20017	
Series D	19225	
Series B	15311	
Series E	12697	
Series F	9932	
Private Equity	7957	
Series H	7244	
Series A	5678	
Series G	3697	
Series J	3370	
Series I	2855	
Seed	1636	
Subsidiary	1094	
NULL	322	

Finally, companies in Post-IPO stage laid off the most during the period.